

Executive Summary

This document will describe the steps done to properly process the data (in this case, we're asked to generate an exponential distribution of about 40 observations). In order to properly analyze the information, 1000 simulations will be run as requested.

After we're done with the above, we have to properly answer the following points:

1. Show where the distribution is centered at and compare it to the theoretical center of the distribution.
2. Show how variable it is and compare it to the theoretical variance of the distribution.
3. Show that the distribution is approximately normal.

Analysis

Let's see the means of each simulation:

```
set.seed(31)
lambda <- 0.2
nsim <- 1000
data <- array(nsim)
for (i in 1 : nsim) { ## 1000 simulations
  data[i] <- data.frame(rexp(40, lambda)) ## 40 observations
}

meanData <- sapply(data, mean) ## take mean of each of the simulations
```

Now, we're going to focus on responding to item 1.

Let's take a look at the distribution of the means (histogram). See Figure 1.

According to the exponential distribution, its theoretical mean is $\frac{1}{\lambda}$, which is 5.0 in the calculation below:

```
1/lambda
```

The distribution of sample means is centered at (which is 4.993867, in the calculation below):

```
mean(meanData)
```

The central limit theorem states that given a distribution with a mean μ and variance σ^2 , the sampling distribution of the mean approaches a normal distribution with a mean μ and $\frac{\sigma^2}{N}$ as variance, when N - the sample size, increases. Both should converge and they did, as expected.

Now, we're going to focus on responding to item 2.

According to the exponential distribution, its theoretical standard deviation is $\frac{1}{\lambda}$, which is 5.0 in the calculation below:

```
1/lambda
```

So, we can conclude that $(\frac{1}{\lambda})^2$ is its theoretical variance, which is 25.0 in the calculation below:

```
(1/lambda)^2
```

The variance of the sampling distribution of the mean is computed as $\frac{\sigma^2}{N}$, where σ is the standard deviation of the distribution and N is the sample size.

So, the variance of the sample means is 0.625, as we can see below:

```
(1/lambda)^2/40
```

If we calculate the actual variance, we can see the value of 0.6291041:

```
var(meanData)
```

The central limit theorem states that given a distribution with a mean μ and variance σ^2 , the sampling distribution of the mean approaches a normal distribution with a mean μ and $\frac{\sigma^2}{N}$ as variance, when N - the sample size, increases. Both should converge and they did, as expected.

Now, we're going to focus on responding the item 3.

Again, the sampling distribution of the mean should approach a normal distribution, according to CLT. Let's check this.

The QQ Plot can help us identify if this is likely to be a normal-like distribution. See Figure 2.

For comparison, let's see a normal distribution with the same parameters as this sample. See Figure 3.

If we put all in one graph, we can see they are very similar. See Figure 4.

The red plot denotes the normal distribution. The other one is our sample.

As the sample size grows, the mean distribution will get more 'normal-like' as stated by CLT.

If we compare the quantiles of our sample with a uniform distribution (as requested by the exercise, in red in graph below), the following will happen (See Figure 5):

Despite the fact that the means and other factors are different, we can see the difference in the graph above: a uniform distribution will have points in both sides of the quantile slope due to its nature of being uniform (and symmetric as we can see) while a normal one will have the points very close to the line.

Appendix

Figure 1:

```
hist(meanData)
```




Figure 2:

```
qqnorm(meanData)
qqline(meanData) ## creates a line that goes from 2nd to the 4th quantile.
```




Figure 3:

```
norm<-rnorm(1000, mean=mean(meanData), sd=sd(meanData))
qqnorm(norm)
qqline(norm) ## creates a line that goes from 2nd to the 4th quantile.
```




Figure 4:

```
q1 <- qqnorm(meanData, plot.it = FALSE)
q2 <- qqnorm(norm, plot.it = FALSE)
plot(range(q1$x, q2$x), range(q1$y, q2$y), type = "n",
      ylab="Sample Quantities", xlab="Theoretical Quantiles")
points(q1)
points(q2, col = "red", pch = 3) ## normal
qqline(meanData) ## quantile line, going from 0.25 to 0.75 quantile
qqline(norm) ## quantile line, going from 0.25 to 0.75 quantile
```




Figure 5:

```
unif<-runif(1000) ## a 1000 uniforms
q1 <- qqnorm(meanData, plot.it = FALSE)
q2 <- qqnorm(unif, plot.it = FALSE)
plot(range(q1$x, q2$x), range(q1$y, q2$y), type = "n",
      ylab="Sample Quantities", xlab="Theoretical Quantiles")
points(q1)
points(q2, col = "red", pch = 3) ## uniform
qqline(meanData) ## quantile line, going from 0.25 to 0.75 quantile
qqline(unif) ## quantile line, going from 0.25 to 0.75 quantile
```

